

Day7 容器进阶之Kubernetes 应用生命周期原理分析

大 纲

- 监控集群组件
- 监控应用
- 管理组件日志
- 管理应用日志
- Deployment升级和回滚
- 配置应用的不同方法
- 应用弹性伸缩
- 应用自恢复

监控集群组件

集群整体状态：

\$ kubectl cluster-info

```
Kubernetes master is running at https://10.142.0.2:6443  
KubeDNS is running at https://10.142.0.2:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
```

更多集群信息：

\$ kubectl cluster-info dump

通过插件部署：

\$ kubectl get pod etcd -n kube-system

\$ kubectl describe pod kube-apiserver -n kube-system

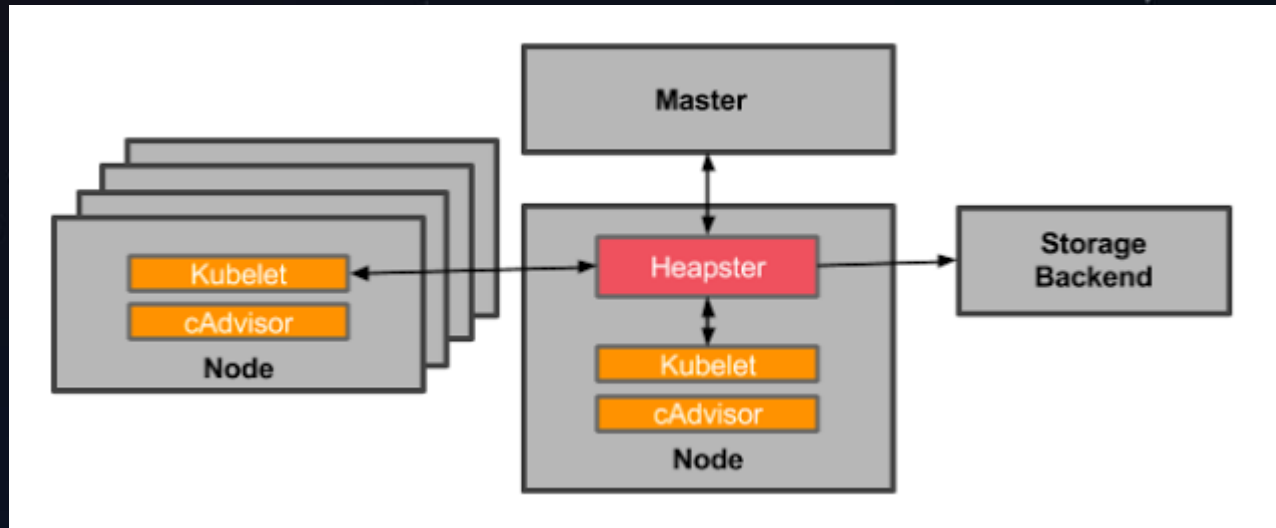
组件metrics：

\$ curl localhost:10250/stats/summary

组件健康状况：

\$ curl localhost:10250/healthz

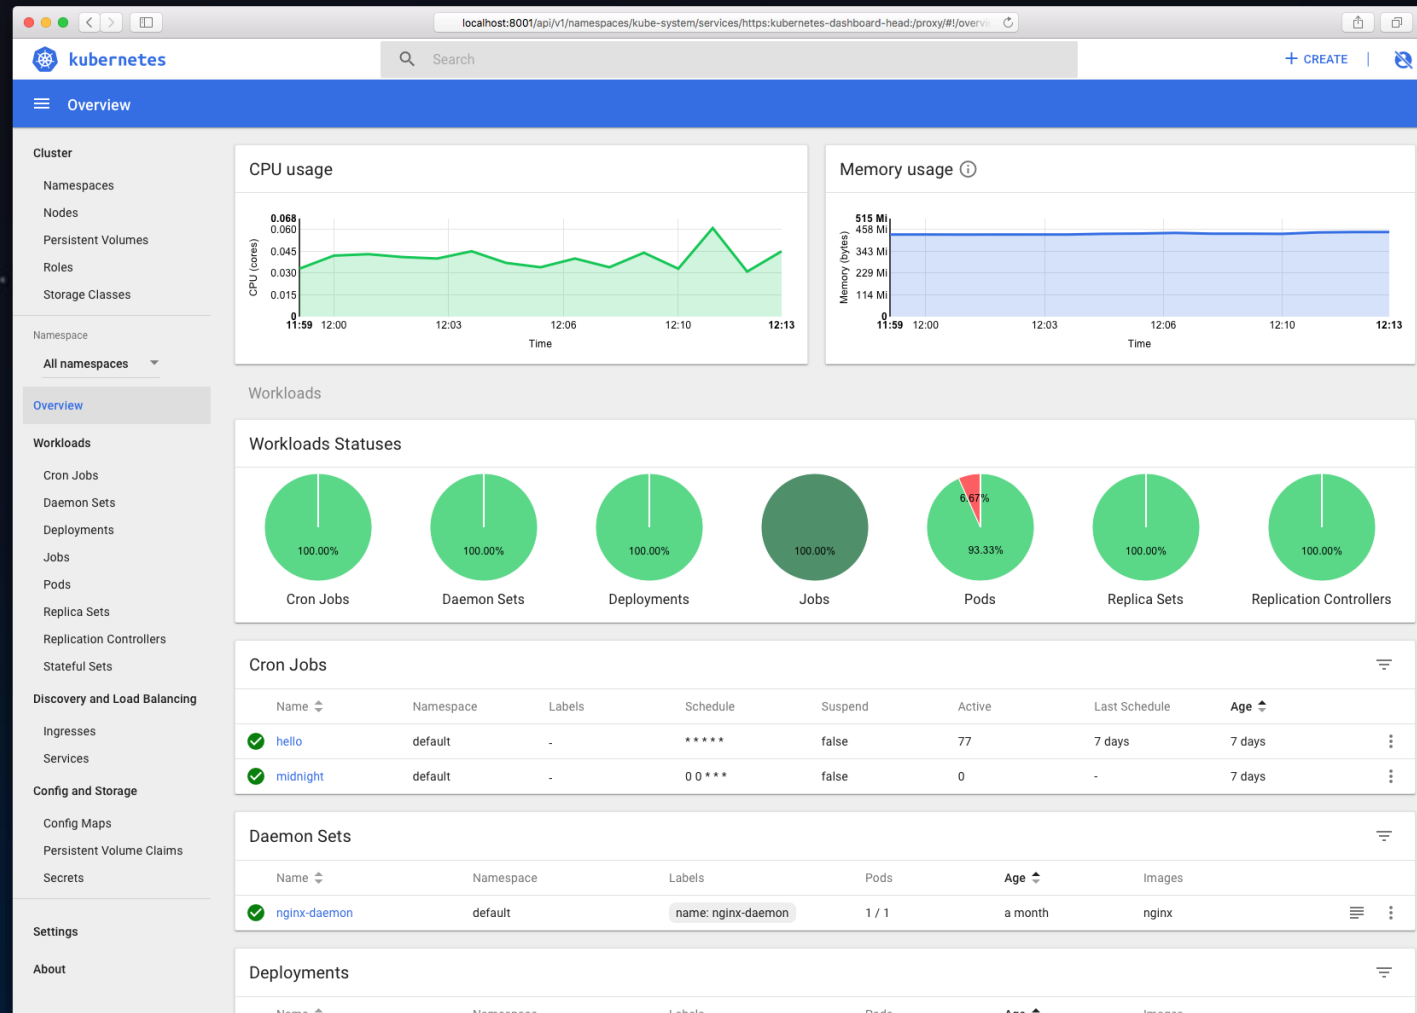
Heapster + cAdvisor监控集群组件



对接了heapster或metrics-server后
展示Node CPU/内存/存储资源消耗:
`$ kubectl top node {node name}`

cAdvisor既能收集容器CPU、内存、文件系统和网络使用统计信息，还能采集节点资源使用情况；
cAdvisor和Heapster都不能进行数据存储、趋势分析和报警。
因此，还需要将数据推送到InfluxDB，Grafana等后端进行存储和图形化展示。
Heapster即将被metrics-server替代

Kuberneetes Dashboard UI



Kubernetes Dashboard用于监控/展示
Kubernetes所有的资源对象：
Cluster (Node , PV等)
Workload (Pod , Deployment等)
Config (Configmap , Secrets等)

...

监控应用

\$ kubectl describe pod

```
Events:
Type      Reason            Age          From          Message
----      -
Normal    Scheduled         1m           default-scheduler    Successfully assigned default/busybox-95444875c-tjcg8 to 127.0.0.1
Normal    SuccessfulMountVolume 1m           kubelet, 127.0.0.1   MountVolume.SetUp succeeded for volume "default-token-8z27r"
Normal    Pulling           24s (x3 over 1m) kubelet, 127.0.0.1   pulling image "busybox"
Warning   Failed            24s (x3 over 1m) kubelet, 127.0.0.1   Failed to pull image "busybox": rpc error: code = Unknown desc = Get https://registry-1.docker.io/v2/: proxyconnect tcp: dial tcp: lookup http on 10.72.55.82:53: no such host
Warning   Failed            24s (x3 over 1m) kubelet, 127.0.0.1   Error: ErrImagePull
Normal    BackOff           10s (x3 over 1m) kubelet, 127.0.0.1   Back-off pulling image "busybox"
Warning   Failed            10s (x3 over 1m) kubelet, 127.0.0.1   Error: ImagePullBackOff
```

对接了heapster或metrics-server后，展示Pod CPU/内存/存储资源消耗：

\$ kubectl top pod {pod name}

```
Every 2.0s: kubectl top pods --namespace backyard

NAME                                CPU(cores)   MEMORY(bytes)
simple-store-mongodb-3006888481-qcp80 5m           324Mi
```

\$ kubectl get pod {pod name} --watch

管理K8S组件日志

组件日志：

/var/log/kube-apiserver.log

/var/log/kube-proxy.log

/var/log/kube-controller-manager.log

/var/log/kubelet.log

使用systemd管理：

\$ journalctl -u kubelet

使用K8S插件部署：

\$ kubectl logs -f kube-proxy

管理K8S应用日志

从容器标准输出截获：

```
$ kubectl logs -f {pod name} -c {container name}
```

```
$ docker logs -f {docker name}
```

日志文件挂载到主机目录：

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: test-pd
```

```
spec:
```

```
  containers:
```

```
  - image: gcr.io/google_containers/test-webserver
```

```
    name: test-container
```

```
    volumeMounts:
```

```
    - mountPath: /log  
      name: log-volume
```

```
  volumes:
```

```
  - name: log-volume
```

```
    hostPath:
```

```
      # directory location on host
```

```
      path: /var/k8s/log
```

直接进入容器内查看日志：

```
$ kubectl exec -it {pod} -c {container} /bin/sh
```

```
$ docker exec -it {container} /bin/sh
```


Deployment升级与回滚 - 1

创建Deployment :

```
$ kubectl run {deployment} --image={image} --replicas={rep.}
```

或使用yaml文件形式，重点配置replicas和image字段。

升级Deployment :

```
$ kubectl set image deployment/nginx-deployment nginx=nginx:1.9.1
```

```
$ kubectl set resources deployment/nginx-deployment -c=nginx --limits=cpu=200m,memory=512Mi
```

升级策略 :

```
minReadySeconds: 5
```

```
strategy:
```

```
  type: RollingUpdate
```

```
  rollingUpdate:
```

```
    maxSurge: 1 #默认25%
```

```
    maxUnavailable: 1 #默认25%
```

Deployment升级与回滚 - 2

暂停Deployment :

```
$ kubectl rollout pause deployment/nginx-deployment
```

恢复Deployment :

```
$ kubectl rollout resume deployment/nginx-deployment
```

查询升级状态 :

```
$ kubectl rollout status deployment/nginx-deployment
```

查询升级历史 :

```
$ kubectl rollout history deploy/nginx-deployment
```

```
$ kubectl rollout history deploy/nginx-deployment --revision=2
```

回滚 :

```
$ kubectl rollout undo deployment/nginx-deployment --to-revision=2
```

应用弹性伸缩

```
$ kubectl scale deployment nginx-deployment --replicas=10
```

对接了heapster，和HPA联动后：

```
$ kubectl autoscale deployment nginx-deployment --min=10 --max=15 --cpu-percent=80
```

应用自恢复 : restartPolicy + livenessProbe

Pod Restart Policy : **Always, OnFailure**, Never

livenessProbe : http/https Get, shell exec, tcpSocket

tcp socket的liveness探针 + always restart例子

apiVersion: v1

kind: Pod

metadata:

name: goproxy

spec:

restartPolicy: Always

containers:

- name: goproxy

image: k8s.gcr.io/goproxy:0.1

ports:

- containerPort: 8080

livenessProbe:

tcpSocket:

port: 8080

initialDelaySeconds: 15

periodSeconds: 20